

Cost & Latency Metrics: Now Live (2026-04-15)

Per-turn cost (USD) and latency (seconds) are now logged as Braintrust **span metrics** on every task row, replacing the earlier attempt to push them through the scorer pipeline (which Braintrust rejected because scores must be in [0,1]).

What changed

Before: `cost_usd_4201` and `latency_ms_4202` were pushed as scorers. `latency_ms_4202` raised `ValueError: score (X) must be between 0 and 1 on every row`; the cost scorer would fail the same way as soon as any turn exceeded \$1.

After: `run_eval.py` calls `current_span().log(metrics={"cost_usd": ..., "latency_seconds": ...})` after each chat turn. The old scorers were deleted from the repo; the Braintrust UI entries should also be deleted to stop them erroring on runs.

Coverage & values — run 2026-04-15 14:50

Metric	Rows covered	Min	Median	Avg	Max	Total
<code>cost_usd</code>	88 / 88	\$0.0024	\$0.124	\$0.109	\$0.274	\$9.62
<code>latency_seconds</code>	88 / 88	1.32 s	30.11 s	28.99 s	69.29 s	—

Both metrics are populated on every task row (no missing values).

Health check on the run

Task-level errors: 0 / 88. **Auth errors:** 0 (JWT refresh logic held across the 30-min run). **Scorer infrastructure errors:** 7 transient timeouts from Braintrust (BrainstoreQuery / 240 s function timeouts), all retried successfully. The experiment-level errors metric dropped from **1.00** in the 14:01 run (old latency scorer failing every row) to **0.08** (only the transient infra timeouts).

Where to see them

In the Braintrust experiment table, `cost_usd` and `latency_seconds` appear next to scorer columns. The experiment diff view highlights regressions when a prompt tweak makes runs more expensive or slower — same surface area as every other scorer, but with accurate numerics (raw USD / seconds, not normalized [0,1]).

Links

Run: [Automated Eval - 2026-04-15 14:50-6e1cf278](#)