

Eval Variance: Two Identical Runs (2026-04-15)

Two full runs of the Benchmark dataset were executed back-to-back against `http://localhost:8001` with the same code (commit `a49bd79`), same dataset (88 rows), same concurrency (3). The only intended difference is LLM nondeterminism. This report quantifies the resulting scorer noise floor.

Runs compared

Baseline: *Automated Eval - 2026-04-15 14:01* (id `45a3da31`)

New: *Automated Eval - 2026-04-15 14:50-6e1cf278* (id `fb030606`)

Per-scorer deltas

Scorer	N	14:01	14:50	Δ (pp)
antisemitism_ab61	37	97.6%	91.9%	-5.7
brand_adherence_98e5	88	47.7%	50.0%	+2.3
concise_and_proportional_81a9	85	93.0%	97.6%	+4.6
delusional_3cd5	12	100.0%	100.0%	+0.0
health_and_judaism_a647	14	86.7%	85.7%	-1.0
html_format_dc7d	76	58.4%	56.6%	-1.9
invalidmissing_reference_scorer_b7ce	50	81.6%	82.0%	+0.4
link_are_valid_06b8	88	88.6%	94.3%	+5.7
minors_88a5	88	95.5%	94.3%	-1.1
non_psak_e2b5	38	100.0%	100.0%	+0.0
politics_7365	88	100.0%	100.0%	+0.0
sefaria_sources_check_5967	71	95.8%	87.3%	-8.5
sefaria_translation_scorer	4	75.0%	25.0%	-50.0
sensitive_jewish_questions_20a7	13	100.0%	100.0%	+0.0
specific_reference_retrieval_e8c4	3	66.7%	66.7%	+0.0
suicide_and_self_harm_7f3b	5	80.0%	100.0%	+20.0
teaser_not_spoiler_7e2c	81	97.6%	97.5%	-0.1
teststandup_ab61	87	100.0%	100.0%	+0.0
theological_questions_ab7a	41	43.2%	56.1%	+12.9
violence_and_hate_speech_754e	46	100.0%	100.0%	+0.0

Highlighted in yellow: scorers with $N \leq 5$ eligible rows. A single flipped LLM-judge verdict moves the percentage by 20–25 pp, so these deltas are dominated by small-sample noise, not real quality change.

What the noise floor looks like

For scorers with $N \geq 40$ (representative sample), 14 of 16 moved within ± 6 pp. The two outliers (`theological_questions` +12.9 pp, `sefaria_sources_check` -8.5 pp) sit on LLM-judge scorers whose per-row verdicts are themselves stochastic — some movement is expected even with identical agent output.

Tiny-N warning

sefaria_translation_scorer (N=4) swung **-50 pp** (75%→25%); suicide_and_self_harm_7f3b (N=5) swung **+20 pp** (80%→100%). With N this small, a single judgment flip is the entire signal. Treat these scorers as directional indicators only — monitor them against bigger samples before drawing conclusions.

Takeaway

Real regressions on this benchmark need to clear roughly ± 6 pp ($N \geq 40$) or ± 25 pp ($N \leq 5$) before they're distinguishable from run-to-run noise. The brand_adherence floor (~50%, stable across both runs) is a consistent quality signal, not variance.