

Eval variance — 2026-04-15, back-to-back runs

Benchmark dataset (n=88), both runs against localhost:8001, concurrency 3. 14:01 = commit a49bd79, 14:50 = commit 05b1f4a. The only code change between them is eval-pipeline plumbing (cost/latency moved from scorer to span metric); agent behavior unchanged. Variance below is LLM nondeterminism.

Scorer	N	14:01	14:50	Δ pp
antisemitism_ab61	37	97.6%	91.9%	-5.7
brand_adherence_98e5	88	47.7%	50.0%	+2.3
concise_and_proportional_81a9	85	93.0%	97.6%	+4.6
delusional_3cd5	12	100.0%	100.0%	+0.0
health_and_judaism_a647	14	86.7%	85.7%	-1.0
html_format_dc7d	76	58.4%	56.6%	-1.9
invalidmissing_reference_scorer_b7ce	50	81.6%	82.0%	+0.4
link_are_valid_06b8	88	88.6%	94.3%	+5.7
minors_88a5	88	95.5%	94.3%	-1.1
non_psak_e2b5	38	100.0%	100.0%	+0.0
politics_7365	88	100.0%	100.0%	+0.0
sefaria_sources_check_5967	71	95.8%	87.3%	-8.5
sefaria_translation_scorer	4	75.0%	25.0%	-50.0
sensitive_jewish_questions_20a7	13	100.0%	100.0%	+0.0
specific_reference_retrieval_e8c4	3	66.7%	66.7%	+0.0
suicide_and_self_harm_7f3b	5	80.0%	100.0%	+20.0
teaser_not_spoiler_7e2c	81	97.6%	97.5%	-0.1
teststandup_ab61	87	100.0%	100.0%	+0.0
theological_questions_ab7a	41	43.2%	56.1%	+12.9
violence_and_hate_speech_754e	46	100.0%	100.0%	+0.0

Yellow rows: $N \leq 5$ — a single flipped LLM-judge verdict shifts these by 20–50 pp, so deltas there are small-sample noise, not quality signal. For $N \geq 40$, 14 of 16 scorers moved within ± 6 pp; `theological_questions` (+12.9) and `sefaria_sources_check` (-8.5) are LLM-judge scorers whose verdicts are themselves stochastic. Working rule: real regression = clearing ± 6 pp ($N \geq 40$) or ± 25 pp ($N \leq 5$).